



# Linear and Logistic Regression

Amos Azaria

# Estimating Galaxy-Phone Cost by Name




# Formalizing Linear Regression

- $y = wx + b$
- What would we do if we had only 2 training examples?  
We could solve 2 equations with 2 parameters
- Our prediction will be  $h(x) = wx + b$
- $Y = \{y_1, y_2, y_3 \dots y_m\}$ ,  $X = \{x_1, x_2, x_3, \dots x_m\}$
- Loss function:  $J(w, b) = \frac{1}{m} \sum_{i=1}^m |wx_i + b - y_i|$
- Or:  $J(w, b) = \frac{1}{2m} \sum_{i=1}^m (wx_i + b - y_i)^2$   
$$J(w, b) = \frac{1}{2m} \sum_{i=1}^m (h(x_i) - y_i)^2$$
- Find  $w, b$  that will minimize  $J(w, b)$

# Gradient Descent

- $N_c$



Schematic of gradient descent.

$w, b$  has  
derivative.

sets of  $f$   
nt update

# What is the gradient ( $\nabla$ )?

- A vector which represents the derivation of a function, which has multiple parameters.
- Each entry is the function's derivative with respect to one of the parameters.

$$f_1(j_1, j_2) = 2j_1 \cdot j_2 + 7j_1$$

$$\nabla(f_1(j_1, j_2)) = (2j_2 + 7, 2j_1)$$

$$f_2(j_1, j_2, j_3) = 3j_1^2 j_2 j_3^3 + 5j_1 j_2$$

$$\nabla(f_2) = (6j_1 j_2 j_3^3 + 5j_2, 3j_1^2 j_3^3 + 5j_1, 9j_1^2 j_2 j_3^2)$$

# The Gradient for Linear Regression

Our loss function:  $J(w, b) = \frac{1}{2m} \sum_{i=1}^m (wx_i + b - y_i)^2$

$$\begin{aligned}\frac{\partial J}{\partial w} &= \frac{1}{2m} \sum_{i=1}^m 2((wx_i + b - y_i)x_i) \\ &= \frac{1}{m} \sum_{i=1}^m (wx_i + b - y_i)x_i\end{aligned}$$

$$\frac{\partial J}{\partial b} = \frac{1}{m} \sum_{i=1}^m (wx_i + b - y_i)$$

$$\nabla(J) = \left( \frac{1}{m} \sum_{i=1}^m (wx_i + b - y_i)x_i, \frac{1}{m} \sum_{i=1}^m (wx_i + b - y_i) \right)$$

# Gradient Descent in Linear Regression

- Pick random  $w$ ,  $b$
- Select the learning rate,  $\alpha$ , (hyper-parameter), e.g. 0.01
- Repeat until convergence:

- Update  $w$  to  $w - \alpha \frac{1}{m} \sum_{i=0}^n x_i (h(x_i) - y_i)$

$$w_{next} = w_{curr} - \alpha \frac{1}{m} \sum_{i=0}^n x_i (w_{curr}(x_i) + b_{curr} - y_i)$$

- Update  $b$  to  $b - \alpha \frac{1}{m} \sum_{i=0}^n 1 \cdot (h(x_i) - y_i)$

$$b_{next} = b_{curr} - \alpha \frac{1}{m} \sum_{i=0}^n 1 \cdot (w_{curr}(x_i) + b_{curr} - y_i)$$

# Linear Regression with SGD in Python

(Using the Galaxy Data-set)

```
import numpy as np
galaxy_data = np.array([[2,70],[3,110],[4,165],[6,390],[7,550]])
w = 0
b = 0
alpha = 0.01
for iteration in range(10000):
    deriv_b = np.mean(1*((w*galaxy_data[:,0]+b)-galaxy_data[:,1]))
    deriv_w = np.mean(galaxy_data[:,0] * ((w*galaxy_data[:,0]+b)-galaxy_data[:,1]))
    b -= alpha*deriv_b
    w -= alpha*deriv_w
    if iteration % 200 == 0 :
        print("it:%d, grad_w:%.3f, grad_b:%.3f, w:%.3f, b:%.3f" %(iteration, deriv_w, deriv_b, w, b))
print("Estimated price for Galaxy S5: ", w*5 + b)
```

$$b_{next} = b_{curr} - \alpha \frac{1}{m} \sum_{i=0}^n 1 \cdot (w_{curr}(x_i) + b_{curr} - y_i)$$

$$w_{next} = w_{curr} - \alpha \frac{1}{m} \sum_{i=0}^n x_i (w_{curr}(x_i) + b_{curr} - y_i)$$



# Linear Regression with SGD in Python

(Using the Galaxy Data-set)

```
import numpy as np
galaxy_data = np.array([[2,70],[3,110],[4,165],[6,390],[7,550]])
w = 0
b = 0
alpha = 0.01
for iteration in range(10000):
    deriv_b = np.mean(1*((w*galaxy_data[:,0]+b)-galaxy_data[:,1]))
    deriv_w = np.mean(galaxy_data[:,0]*((w*galaxy_data[:,0]+b)-galaxy_data[:,1]))
    b -= alpha*deriv_b
    w -= alpha*deriv_w
    if iteration % 200 == 0 :
        print("it:%d, grad_w:%.3f, grad_b:%.3f, w:%.3f, b:%.3f" %(iteration, deriv_w, deriv_b, w, b))
print("Estimated price for Galaxy S5: ", w*5 + b)
```

$$b_{next} = b_{curr} - \alpha \frac{1}{m} \sum_{i=0}^n 1 \cdot (w_{curr}(x_i) + b_{curr} - y_i)$$

$$w_{next} = w_{curr} - \alpha \frac{1}{m} \sum_{i=0}^n x_i (w_{curr}(x_i) + b_{curr} - y_i)$$

# Value of b in a certain iteration

```
galaxy_data = np.array([[2,70],[3,110],[4,165],[6,390],[7,550]])  
deriv_b = np.mean(1*((w*galaxy_data[:,0]+b)-galaxy_data[:,1]))
```

w=11 ; b=5

$$\text{mean} \left( 11 \times \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix} + 5 - \begin{bmatrix} 70 \\ 110 \\ 165 \\ 390 \\ 550 \end{bmatrix} \right) = \text{mean} \begin{bmatrix} 22 + 5 - 70 \\ 33 + 5 - 110 \\ 44 + 5 - 165 \\ 66 + 5 - 390 \\ 77 + 5 - 550 \end{bmatrix} = \text{mean} \begin{bmatrix} -43 \\ -72 \\ -116 \\ -319 \\ -468 \end{bmatrix} = -203.6$$

```
b -= alpha*deriv_b
```

```
b = 5 - 0.1*(-203.6) = 25.36
```

# Value of w in a certain iteration

```
galaxy_data = np.array([[2,70],[3,110],[4,165],[6,390],[7,550]])  
deriv_w = np.mean(galaxy_data[:,0]*((w*galaxy_data[:,0]+b)-galaxy_data[:,1]))
```

w=11 ; b=5

$$\text{mean} \left( \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix} \times \left( 11 \times \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix} + 5 - \begin{bmatrix} 70 \\ 110 \\ 165 \\ 390 \\ 550 \end{bmatrix} \right) \right) = \text{mean} \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix} \times \begin{bmatrix} 11 * 2 + 5 - 70 \\ 11 * 3 + 5 - 110 \\ 11 * 4 + 5 - 165 \\ 11 * 6 + 5 - 390 \\ 11 * 7 + 5 - 550 \end{bmatrix} = \text{mean} \left( \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix} \times \begin{bmatrix} -43 \\ -72 \\ -116 \\ -319 \\ -468 \end{bmatrix} \right) = \text{mean} \begin{bmatrix} -86 \\ -219 \\ -464 \\ -357 \\ -3276 \end{bmatrix} = -880.4$$

w -= alpha\*deriv\_w

w = 11 - 0.1\*(-880.4) = 99.04

# Result:

it:0, grad\_w:1464.000, grad\_b:257.000, w:14.640, b:2.570  
it:200, grad\_w:3.825, grad\_b:-19.693, w:70.598, b:-33.968  
it:400, grad\_w:2.859, grad\_b:-14.720, w:77.230, b:-68.115  
it:600, grad\_w:2.137, grad\_b:-11.003, w:82.188, b:-93.639  
it:800, grad\_w:1.597, grad\_b:-8.224, w:85.893, b:-112.717  
it:1000, grad\_w:1.194, grad\_b:-6.147, w:88.663, b:-126.977  
it:1400, grad\_w:0.667, grad\_b:-3.434, w:92.280, b:-145.604  
it:1800, grad\_w:0.373, grad\_b:-1.919, w:94.301, b:-156.010  
it:2600, grad\_w:0.116, grad\_b:-0.599, w:96.062, b:-165.073  
it:3600, grad\_w:0.027, grad\_b:-0.140, w:96.674, b:-168.226  
it:4400, grad\_w:0.008, grad\_b:-0.044, w:96.802, b:-168.886  
it:5400, grad\_w:0.002, grad\_b:-0.010, w:96.847, b:-169.116  
it:6200, grad\_w:0.001, grad\_b:-0.003, w:96.856, b:-169.164  
it:7000, grad\_w:0.000, grad\_b:-0.001, w:96.859, b:-169.179  
it:7800, grad\_w:0.000, grad\_b:-0.000, w:96.860, b:-169.184  
it:9800, grad\_w:0.000, grad\_b:-0.000, w:96.860, b:-169.186

Estimated price for Galaxy S5: 315.116281573

- Actual price was: **\$250**
- Estimated price for Galaxy S1 is: **-72.22** ☹️
- What would happen with  $\alpha=0.5$ ?

# Adding Features

- Screen size
- Number of cores
- Core speed
- ...

# Adding Features (cont.)

- $x_1 = \{x_{11}, x_{12}, x_{13}, \dots, x_{1k}\}$
- $W = \{w_1, w_2, w_3, \dots, w_k\}$
- $\text{Loss}(W, b) = \frac{1}{2m} \sum (y - (Wx + b))^2$
- Sometimes we set:  $x_{i0} = 1$  ;  $w_0 = b$   
So that  $x_i = \{x_{i0}, x_{i1}, x_{i2}, x_{i3}, \dots, x_{ik}\}$   
=> no need for “b”
- $W = \{w_0 = b, w_1, w_2, \dots\}$
- $\text{Loss} = \frac{1}{2m} \sum (y - Wx)^2$

# Gradient Descent with Multiple Features (is actually the same...)

- $\nabla$  is:  $\sum_{i=0}^n \overleftarrow{x}_i (h(\overleftarrow{x}_i) - y_i)$
  - Initialize  $W$
  - Repeat:
    - Calculate  $\nabla: \frac{1}{m} \sum_{i=0}^n \overleftarrow{x}_i (h(\overleftarrow{x}_i) - y_i)$
- Update:  $W = W - \alpha \nabla$

# Adding quadratic feature

- We can add a feature which is simply a square of the first feature.
- We will end up with 2 weights and one bias:
  - $w_1 x + w_2 x^2 + b$
- This way, instead of fitting a linear line, we are actually fitting a quadratic function (parabola).



# Quadratic Feature

```
import numpy as np
data_x = np.array([[2,4],[3,9],[4,16],[6,36],[7,49]])
data_y = np.array([70,110,165,390,550])
w1 = 0
w2 = 0
b = 0
alpha = 0.001
for iteration in range(1000000):
    deriv_b = np.mean(1*(w1*data_x[:,0]+w2*data_x[:,1]+b)-data_y))
    deriv_w1 = np.dot(((w1*data_x[:,0]+w2*data_x[:,1]+b)-data_y), data_x[:,0]) *
1.0/len(data_y)
    deriv_w2 = np.dot(((w1*data_x[:,0]+w2*data_x[:,1]+b)-data_y), data_x[:,1]) *
1.0/len(data_y)
    b -= alpha * deriv_b
    w1 -= alpha * deriv_w1
    w2 -= alpha * deriv_w2

print("Estimated price for Galaxy S5: ", np.dot(np.array([5,25]),np.array([w1, w2])) + b)
print("Estimated price for Galaxy S1: ", np.dot(np.array([1,1]),np.array([w1, w2])) + b)
```

$h(x) - y$

# Value of b in a certain iteration

```
data_x = np.array([[2,4],[3,9],[4,16],[6,36],[7,49]])
data_y = np.array([70,110,165,390,550])
deriv_b = np.mean(1*((w1*data_x[:,0]+w2*data_x[:,1]+b)-data_y))
```

w1=11; w2=7 ; b=5

$$\text{mean} \left( 11 \times \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix} + 7 \times \begin{bmatrix} 4 \\ 9 \\ 16 \\ 36 \\ 49 \end{bmatrix} + 5 - \begin{bmatrix} 70 \\ 110 \\ 165 \\ 390 \\ 550 \end{bmatrix} \right) = \text{mean} \begin{bmatrix} 22 + 28 + 5 - 70 \\ 33 + 63 + 5 - 110 \\ \dots \\ \dots \\ \dots \end{bmatrix} = \text{mean} \begin{bmatrix} -15 \\ -9 \\ \dots \\ \dots \\ \dots \end{bmatrix} = \dots$$

b -= alpha \* deriv\_b

# Value of w1 in a certain iteration

```
data_x = np.array([[2,4],[3,9],[4,16],[6,36],[7,49]])
data_y = np.array([70,110,165,390,550])
deriv_w1 = np.dot(((w1*data_x[:,0]+w2*data_x[:,1]+b)-data_y), data_x[:,0])* 1.0/len(data_y)
```

w1=11; w2=7 ; b=5

$$\left( 11 \times \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix} + 7 \times \begin{bmatrix} 4 \\ 9 \\ 16 \\ 36 \\ 49 \end{bmatrix} + 5 - \begin{bmatrix} 70 \\ 110 \\ 165 \\ 390 \\ 550 \end{bmatrix} \right) \times \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix} = \begin{bmatrix} 22 + 28 + 5 - 70 \\ 33 + 63 + 5 - 110 \\ \dots \\ \dots \\ \dots \end{bmatrix} \times \begin{bmatrix} 2 \\ 3 \\ 4 \\ 6 \\ 7 \end{bmatrix}$$


---


$$\frac{\text{len}(\text{data\_y})}{5} = \frac{-1350}{5}$$

```
w1 -= alpha * deriv_w1
w1 = 0.001 * (-270) = -0.27
```

# Quadratic Feature (Weights as Vectors)

```
data_x = np.array([[2,4],[3,9],[4,16],[6,36],[7,49]])
```

```
data_y = np.array([70,110,165,390,550])
```

```
w = np.array([0.,0])
```

```
b = 0
```

```
alpha = 0.001
```

```
for iteration in range(1000000):
```

```
    deriv_b = np.mean(1*((np.dot(data_x,w)+b)-data_y))
```

```
    gradient_w = 1.0/len(data_y) * np.dot(((np.dot(data_x,w)+b)-data_y), data_x)
```

```
    b -= alpha*deriv_b
```

```
    w -= alpha*gradient_w
```

```
print("Estimated price for Galaxy S5: ", np.dot(np.array([5,25]),w) + b)
```

```
print("Estimated price for Galaxy S1: ", np.dot(np.array([1,1]),w) + b)
```

$$\nabla_{\vec{w}} = \left( \begin{bmatrix} 2 & 4 \\ 3 & 9 \\ 4 & 16 \\ 6 & 36 \\ 7 & 49 \end{bmatrix} \times \begin{bmatrix} w_1 \\ w_2 \end{bmatrix} + b - \begin{bmatrix} 70 \\ 110 \\ 165 \\ 390 \\ 550 \end{bmatrix} \right) \times \begin{bmatrix} 2 & 4 \\ 3 & 9 \\ 4 & 16 \\ 6 & 36 \\ 7 & 49 \end{bmatrix} = \begin{bmatrix} 2 * w_1 + 4 * w_2 + b - 70 \\ \dots \\ \dots \\ \dots \\ \dots \end{bmatrix} \times \begin{bmatrix} 2 & 4 \\ 3 & 9 \\ 4 & 16 \\ 6 & 36 \\ 7 & 49 \end{bmatrix} = \begin{bmatrix} \nabla w_1 \\ \nabla w_2 \end{bmatrix}$$

# Results

Prediction for Galaxy S5: \$263.63

(actual: \$250) 😊

Prediction for Galaxy S1: \$71.81


(actual: \$30) 😞

# Adding (too) Many Features

- We can add more and more features by adding additional polynomial terms (e.g.  $x^3$ ,  $x^4$ , etc.)
- (We need to make sure not to overflow, so we should really normalize the values we get).
- Should we add 20 features? What can happen?

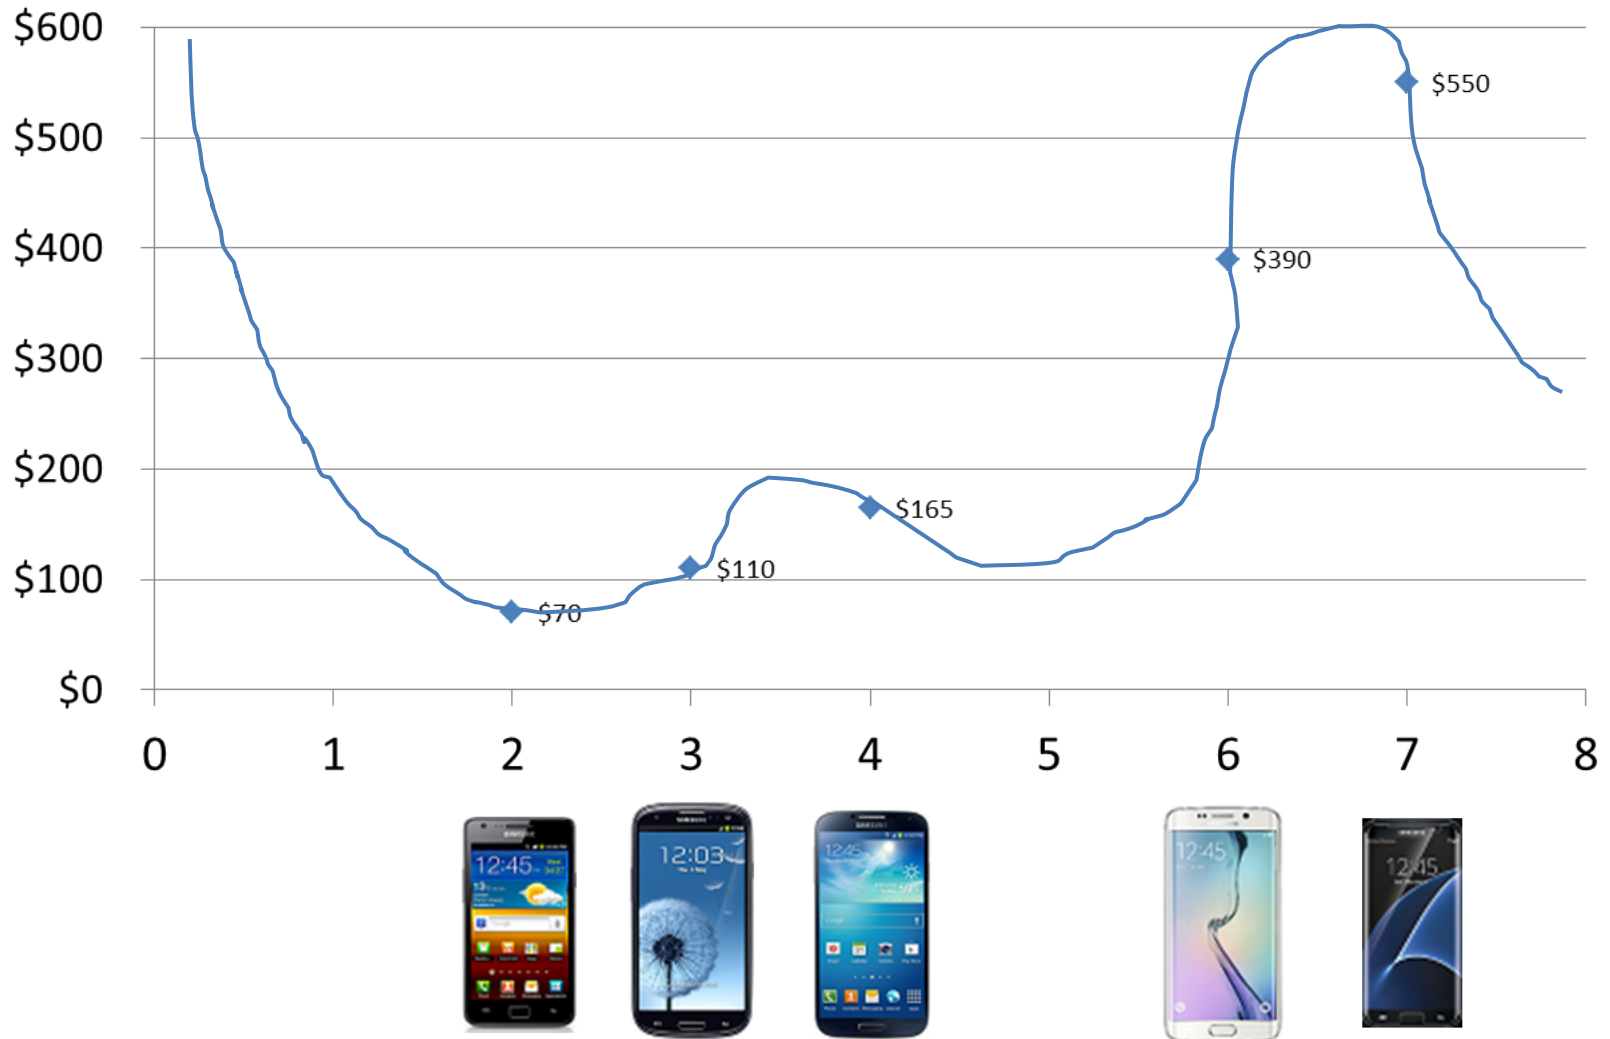
# Over-Fitting

Loss =  $5.3 \times 10^{-7}$ !!!



# Over-Fitting


Galaxy-Price on ebay






# Results with "only" 10 features

Loss = 0.0049



# Train/Test error



Credit: Collindcarroll.com

# Train-Test-Validation

- If we want the test set to be a good prediction to what will happen with new data, it should be used only once.
- Therefore, we may want to have a validation set.
- The hyper-parameters / model complexity will be determined such that they maximize the accuracy of the validation set.

# Closed form solution

- Linear regression has a closed-form solution:

$$W = (X^T X)^{-1} X^T Y$$

- We won't be focusing on it since we will be moving beyond linear regression to problems that do not have closed-form solution.
- Furthermore, the closed form solution may be less practical for big data.
- We will be using gradient descent (and its variants) until the end of the course.

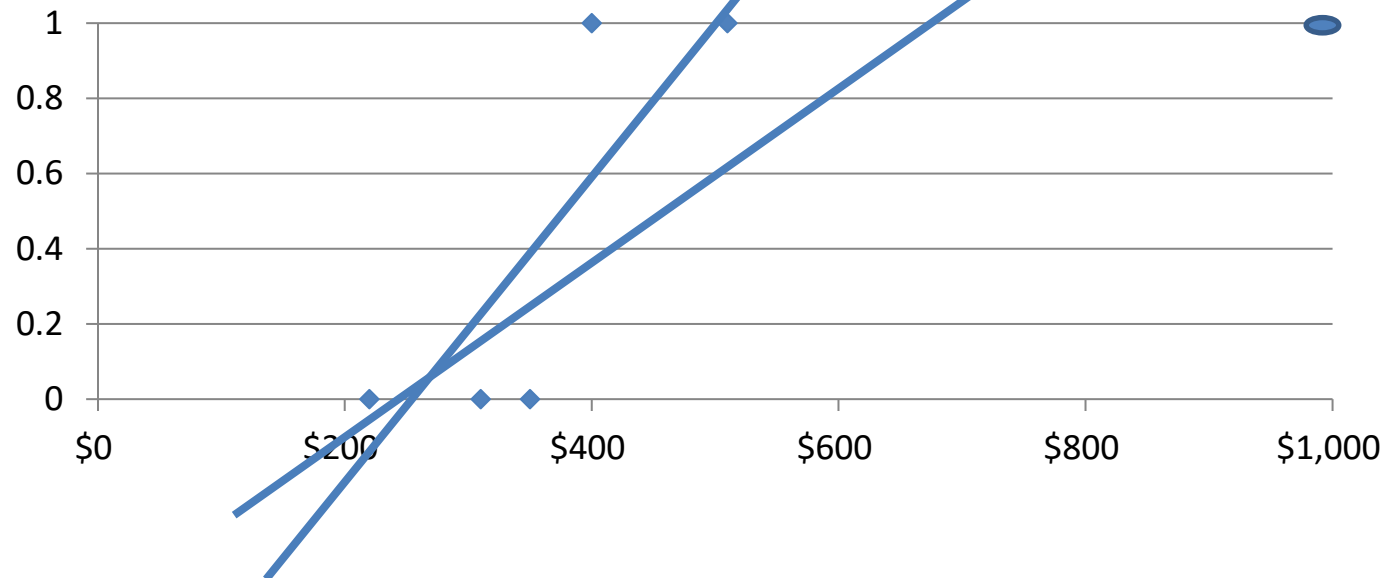
# BGD, SGD, MB-GD

- Batch Gradient Descent (BGD): uses **all Data-set** to compute gradient (this is what we learnt).
  - May be too large, or take too long:
- Stochastic Gradient Descent (SGD): uses **only a single example** at a time (shuffle data first):
  - More iterations, since each iteration is less accurate
- Mini-Batch Gradient Descent (MB-GD): uses only **a subset** of the data-set, (e.g. 50) at a time:
  - A compromise which takes advantage of vectorization.

**CLASSIFICATION**

# Classification


- Suppose we wanted to classify phones into new (1) or old/used (0), using the price as a single feature.
- Could we use linear regression?



# Logistic Regression

- The Logistic function is:

- $h(x) = \frac{1}{1 + e^{-(Wx+b)}}$



- A prediction of 1 will mean that we are **certain** that the value is 1.
- Instead of using least squares, we will use the following loss function:

This is actually the loss function we get when we try to maximize the likelihood of the data

$$J(w, b) = -\frac{1}{m} \sum_{i=1}^m (y_i(\log(h(x_i))) + (1 - y_i)\log(1 - h(x_i)))$$

- It turns out that the Gradient of the loss is:


- $\frac{1}{m} \sum_{i=0}^n x_i (h(x_i) - y_i)$



# Logistic Regression

- The Logistic function is:

- $h(x) = \frac{1}{1 + e^{-(Wx+b)}}$



- A prediction of 1 will mean that we are **certain** that the value is 1.
- Instead of using least squares, we will use the following loss function:

This is actually the loss function we get when we try to maximize the likelihood of the data

$$J(w, b) = -\frac{1}{m} \sum_{i=1}^m (y_i(\log(h(x_i))) + (1 - y_i)\log(1 - h(x_i)))$$

- It turns out that the Gradient of the loss is:

- $\frac{1}{m} \sum_{i=0}^n x_i (h(x_i) - y_i)$

# Gradient Descent in Logistic Regression

- Pick random  $w, b$
- Select the learning rate,  $\alpha$ , (hyper-parameter), e.g. 0.01
- Repeat until convergence:
  - Update  $w$  to  $w - \alpha \frac{1}{m} \sum_{i=0}^n x_i (h(x_i) - y_i)$
  - Update  $b$  to  $b - \alpha \frac{1}{m} \sum_{i=0}^n 1 \cdot (h(x_i) - y_i)$

# How-come?

- Does the previous slide look familiar?
- The previous slide is identical to the slide we have seen in linear regression (I actually did a copy-paste and only changed the title).
- So how is logistic regression actually different than linear regression? If it is exactly the same algorithm, why not use linear regression?

----

Obviously, the algorithms are different because the hypothesis ( $h(x)$ ) is totally different...

# Employed or not?

- We have a data-base with all our users and we want to send out job-offers.
- For that, we need to know all unemployed users, though we only know this information on a fraction of the data.
- We would like to build a classifier that determines whether a user is employed or not, based on the user's age, gender and years of experience.

# Our dataset

This is fake data, sorry about any gender/age biases introduced intentionally...

- Employed users:
  - Female, 28 years old, 4 years of experience
  - Female, 60 years old, 34 years of experience
  - Female, 25 years old, 3 year of experience
  - Male, 54 years old, 20 years of experience
  - Male, 24 years old, 2 years of experience
  - Male, 39 years old, 12 years of experience
  - Male, 30 years old, 4 years of experience
- Unemployed users:
  - Female, 36 years old 10 years of experience
  - Female, 26 years old 1 year of experience
  - Male, 44 years old, 9 years of experience

Do you think that a female, 49 year old with 8 years of experience employed?

What about a male, 29 years old with 3 years of experience?

And if it were a female?

```
import numpy as np
data_x = np.array([[1,28,4],[1,60,34],[1,25,3],[0,54,20],[0,24,2],[0,39,12],[0,30,4],[1,3
data_y = np.array([1,1,1,1,1,1,1,0,0,0])
```

```
def h(x,w,b):
    return 1 / (1+np.exp(-(np.dot(x,w) + b)))
```

$$h(x) = \frac{1}{1 + e^{-(Wx+b)}}$$

```
w = np.array([0.,0,0])
```

```
b = 0
```

```
alpha = 0.001
```

```
for iteration in range(100000):
```

```
    gradient_b = np.mean(1*(data_y-(h(data_x,w,b))))
```

```
    gradient_w = np.dot((data_y-h(data_x,w,b)), data_x)*1/len(data_y)
```

```
    b += alpha*gradient_b
```

```
    w += alpha*gradient_w
```

```
print("User [1, 49, 8] prob of working: ", h(np.array([[1, 49, 8]]),w,b))
```

```
print("User [0, 29, 3] prob of working: ", h(np.array([[0, 29, 3]]),w,b))
```

```
print("User [1, 29, 3] prob of working: ", h(np.array([[1, 29, 3]]),w,b))
```

# Results

User [1, 49, 8] prob of working: [ 0.21107079]

User [0, 29, 3] prob of working: [ 0.66430518]

User [1, 29, 3] prob of working: [ 0.43735087]

# Meaning of Logistic Regression Result

- The result given by logistic regression is suppose to relate to the probability of the instance belonging to the class  $p(y=1 \mid X)$ .
- Logistic regression is a discriminative model, that is, it tries to model  $p(y \mid X)$  directly.
- Remember that in the naïve Bayes classifier (which is a generative model), we used the Bayes rule, and therefore had to model:

$p(y)$  and  $p(X \mid y)$  (and  $p(X)$ )

$$p(Y \mid X) = \frac{p(Y)p(X|Y)}{p(X)}$$



# Multiple Classes

- Many times classification is into several labels. E.g. Classify an image to an object:  
cat/dog/airplane/sea/house
- We could build 5 different classifiers...
- More often we use one-hot vector representations for the labels. E.g.: cat = [1, 0, 0, 0, 0], sea = [0, 0, 0, 1, 0]
- **SoftMax**: as if we do logistic regression for each output alone and then scale:

$$\bullet h(y = i | x) = \frac{e^{(W_i x + b_i)}}{\sum_j^k e^{(W_j x + b_j)}}$$

– Note that:  $\sum_{i \in \{0, 1, \dots, k-1\}} h(y = i | x) = 1$



# Closing notes

**3 DATABASE ADMINS  
WALKED INTO  
A NOSQL BAR...**

**A LITTLE WHILE LATER  
THEY WALKED OUT BECAUSE  
THEY COULDN'T FIND A TABLE**

# Closing notes

- Data and Database
- Sql databases
- Normalization
- XML & Json
- NoSql databases
- Java Streams & Spark
- Introduction to Machine Learning

**THANK YOU FOR**

**LISTENING**

